

mmap — Memory-map Files

Purpose: Memory-map files instead of reading the contents directly.

Memory-mapping a file uses the operating system virtual memory system to access the data on the file system directly, instead of using normal I/O functions. Memory-mapping typically improves I/O performance because it does not involve a separate system call for each access and it does not require copying data between buffers - the memory is accessed directly by both the kernel and the user application.

Memory-mapped files can be treated as mutable strings or file-like objects, depending on the need. A mapped file supports the expected file API methods, such as close(), flush(), read(), readline(), seek(), tell(), and write(). It also supports the string API, with features such as slicing and methods like find().

All of the examples use the text file lorem.txt, containing a bit of Lorem Ipsum. For reference, the text of the file is

lorem.txt

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Donec egestas, enim et consectetuer ullamcorper, lectus ligula rutrum leo, a elementum elit tortor eu quam. Duis tincidunt nisi ut ante. Nulla facilisi. Sed tristique eros eu libero. Pellentesque vel arcu. Vivamus purus orci, iaculis ac, suscipit sit amet, pulvinar eu, lacus. Praesent placerat tortor sed nisl. Nunc blandit diam egestas dui. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Aliquam viverra fringilla leo. Nulla feugiat augue eleifend nulla. Vivamus mauris. Vivamus sed mauris in nibh placerat egestas. Suspendisse potenti. Mauris massa. Ut eget velit auctor tortor blandit sollicitudin. Suspendisse imperdiet justo.

Note

There are differences in the arguments and behaviors for mmap() between Unix and Windows, which are not fully discussed here. For more details, refer to the standard library documentation.

Reading

Use the mmap() function to create a memory-mapped file. The first argument is a file descriptor, either from the fileno() method of a file object or from os.open(). The caller is responsible for opening the file before invoking mmap(), and closing it after it is no longer needed.

The second argument to mmap() is a size in bytes for the portion of the file to map. If the value is 0, the entire file is mapped. If the size is larger than the current size of the file, the file is extended.

Note

Windows does not support creating a zero-length mapping.

An optional keyword argument, access, is supported by both platforms. Use ACCESS READ for read-only access, ACCESS WRITE for write-through (assignments to the memory go directly to the file), or ACCESS COPY for copy-on-write (assignments to memory are not written to the file).

```
# mmap_read.py
import mmap
with open('lorem.txt', 'r') as f:
    with mmap.mmap(f.fileno(), 0,
                   access=mmap.ACCESS READ) as m:
        print('First 10 bytes via read :', m.read(10))
```

The file pointer tracks the last byte accessed through a slice operation. In this example, the pointer moves ahead 10 bytes after the first read. It is then reset to the beginning of the file by the slice operation, and moved ahead 10 bytes again by the slice. After the slice operation, calling read() again gives the bytes 11-20 in the file.

```
$ python3 mmap_read.py
First 10 bytes via read : b'Lorem ipsu'
First 10 bytes via slice: b'Lorem ipsu'
2nd 10 bytes via read : b'm dolor si'
```

Writing

To set up the memory mapped file to receive updates, start by opening it for appending with mode 'r+' (not 'w') before mapping it. Then use any of the API methods that change the data (write(), assignment to a slice, etc.).

The next example uses the default access mode of ACCESS WRITE and assigning to a slice to modify part of a line in place.

```
# mmap write slice.py
import mmap
import shutil
# Copy the example file
shutil.copyfile('lorem.txt', 'lorem_copy.txt')
word = b'consectetuer'
reversed = word[::-1]
print('Looking for
                    :', word)
print('Replacing with :', reversed)
with open('lorem copy.txt', 'r+') as f:
    with mmap.mmap(f.fileno(), 0) as m:
        print('Before:\n{}'.format(m.readline().rstrip()))
       m.seek(0) # rewind
        loc = m.find(word)
        m[loc:loc + len(word)] = reversed
        m.flush()
        m.seek(0) # rewind
        print('After :\n{}'.format(m.readline().rstrip()))
        f.seek(0) # rewind
        print('File :\n{}'.format(f.readline().rstrip()))
```

The word "consectetuer" is replaced in the middle of the first line in memory and in the file.

```
$ python3 mmap_write_slice.py

Looking for : b'consectetuer'
Replacing with : b'reutetcesnoc'
Before:
b'Lorem ipsum dolor sit amet, consectetuer adipiscing elit.'
After :
b'Lorem ipsum dolor sit amet, reutetcesnoc adipiscing elit.'
File :
Lorem ipsum dolor sit amet, reutetcesnoc adipiscing elit.'
```

Copy Mode

Using the access setting ACCESS_COPY does not write changes to the file on disk.

```
# mmap_write_copy.py
import mmap
import mmap
```

```
IMPORT SHUTIL
# Copy the example file
shutil.copyfile('lorem.txt', 'lorem copy.txt')
word = b'consectetuer'
reversed = word[::-1]
with open('lorem_copy.txt', 'r+') as f:
    with mmap.mmap(f.fileno(), 0,
                   access=mmap.ACCESS COPY) as m:
        print('Memory Before:\n{}'.format(
            m.readline().rstrip()))
        print('File Before :\n{}\n'.format(
            f.readline().rstrip()))
        m.seek(0) # rewind
        loc = m.find(word)
        m[loc:loc + len(word)] = reversed
        m.seek(0) # rewind
        print('Memory After :\n{}'.format(
            m.readline().rstrip()))
        f.seek(0)
        print('File After
                          :\n{}'.format(
            f.readline().rstrip()))
```

It is necessary to rewind the file handle in this example separately from the mmap handle, because the internal state of the two objects is maintained separately.

```
$ python3 mmap_write_copy.py

Memory Before:
b'Lorem ipsum dolor sit amet, consectetuer adipiscing elit.'
File Before :
Lorem ipsum dolor sit amet, consectetuer adipiscing elit.'

Memory After :
b'Lorem ipsum dolor sit amet, reutetcesnoc adipiscing elit.'
File After :
Lorem ipsum dolor sit amet, consectetuer adipiscing elit.'
```

Regular Expressions

Since a memory mapped file can act like a string, it can be used with other modules that operate on strings, such as regular expressions. This example finds all of the sentences with "nulla" in them.

Because the pattern includes two groups, the return value from findall() is a sequence of tuples. The print statement pulls out the matching sentence and replaces newlines with spaces so each result prints on a single line.

```
$ python3 mmap_regex.py
b'Nulla facilisi.'
b'Nulla feugiat augue eleifend nulla.'
```

See also

- Standard library documentation for mmap
- Python 2 to 3 porting notes for mmap
- os The os module.
- <u>re</u> Regular expressions.

◆ filecmp — Compare Files

codecs — String Encoding and Decoding €

Quick Links

Reading Writing Copy Mode **Regular Expressions**

This page was last updated 2016-12-28.

Navigation

filecmp — Compare Files

codecs — String Encoding and Decoding



Get the book

The output from all the example programs from PyMOTW-3 has been generated with Python 3.7.1, unless otherwise noted. Some of the features described here may not be available in earlier versions of Python.

Looking for examples for Python 2?

This Site

■ Module Index

 \boldsymbol{I} Index









© Copyright 2019, Doug Hellmann



Other Writing



The Python Standard Library By Example